

Temporal Planning and Inferencing for Personal Task Management with SPSE2

Andrew Dougherty

FRDCSA Project

1125 Village Center Pkwy Unit 1. Aurora, IL 60506

andrewdo@frdcsa.org

Abstract

SPSE2 is a free and open source software system for personal task management that uses publicly available temporal planning systems and tools. It provides a GUI for users to edit planning domains consisting of assertions about goals and their interrelations and temporal, cost and other constraints. It also integrates calendar and transportation planning information. A voice-controlled mobile application on the Android platform for goal setting, interactive execution monitoring and replanning is currently under development. The rest of this paper concerns details and planned features, application areas and future work; as well as various other AI-related technologies being used to extend the usefulness of the core functionality. For example, computational semantics software will extract models of the meaning of the goals and translate into elements of the planning language, and deontic logics will help to determine the consistency of the goals and plans with the user's preferred ethics and value systems. The work is being conducted as part of the Cognitive Prostheses research focus of the FRDCSA project.

Introduction

This paper documents progress within the Formalized Research Database; Cluster, Study and Apply (FRDCSA) project towards applying existing AI planning and scheduling tools to the domain of personal task management. This tool is being developed for general use, but particularly for 1) those with disabilities impairing so-called executive skills, and 2) those in poverty who can only secure basic needs with difficulty. Disabilities that affect executive skills include Pervasive Developmental Disorders such as Autism and Schizophrenia. "They are called executive skills because they help people execute tasks. Every person has a set of 12 executive skills (self-restraint, working memory, emotion control, focus, task initiation, planning/prioritization, organization, time management, defining and achieving goals, flexibility, observation and stress tolerance)" (Martin 2006). Additionally, the proliferation of inexpensive Wi-Fi-enabled, Android-based smart phones, combined with free software and perhaps solar-power chargers, has provided a novel vector to help to supply the

homeless or similarly disadvantaged with practical, survival-oriented computing resources. Anecdotal evidence suggests further that the executive skills of those even without such limiting factors still may have room for improvement, especially when compared to a more optimal automated negotiation and coordination of temporal constraints and privacy preferences. Such improvements in local planning efficacy and resource utilization could have macroscopic effects and relax otherwise exacerbated resource conflicts. Additionally, the application of user-defined ethical analytics through computational deontology offers a method to constrain the space of possible actions that are available to intelligent personal agents. Therefore, development of a free and open source (FOSS) personal task management system that interfaces with existing sources of data such as calendars, routing applications and mobile phones promises to help to improve the quality of life generally.

Overview of the FRDCSA Project

The first and major goal of the FRDCSA project, 11 years and running, is to help to provide for better security and quality of life for all sentient beings. A major assumption is that FOSS artificial intelligence, engineered correctly and with unlimited redistribution, satisfies this goal. To avoid polemics, the project is concerned only with implementing a restricted form of weak AI. The approach, motivated by algorithmic information theory and information-theoretic computational complexity of metamathematics (Chaitin 1974), has two prongs - to develop an increasingly complete theorem proving system and library (called Formalized Research Database (FRD)), and to develop an increasingly complete collection of practical software (Cluster, Study and Apply (CSA)). Perhaps, the two approaches are in some sense theoretically equivalent by the Curry-Howard isomorphism, which would seem to say that the classes of programs and proofs are coextensive. So ideally the FRD is a practical, transfinite implementation of Hilbert's program, which can never be completed, but which can, by engineering a sequence of logics each more complete than the previous, decide in the limit all problems that are not absolutely undecidable (should any of those exist). This idea may be similar to or the same as one of Turing of creating a sequence of logics, each more complete than the previous, based upon the assumption of the existence of in-

creasingly large constructible ordinals (Turing 1939). Alternatively, the CSA consists of a set of programs for building and packaging most or all known freely available software systems and datasets.

The primary goal of the project and limitations to personal productivity necessitate the development of tools, including a personal task management system, which provide tangible benefits to all users and in particular those working on FOSS software development. Yet such a system has been reasonably difficult to engineer in the context of the project, as evidenced by the almost 10 such systems in various stages of completion, many not mentioned here. Also, it has been difficult to find an accurate name for this entire collection of functionality.

Overview of the SPSE2 System

SPSE2 is a system for personal task management. It provides a GUI for users to edit planning domains. These domains consist of assertions in a knowledge base regarding goals, their interrelations, and constraints of types such as temporal and cost. This GUI allows users to specify the constraints, and then the domain is converted into a PDDL domain and a plan generated. Future versions will include a fully functional and re-engineered interactive execution monitor that will walk the users through the plans and allow replanning from hands-free, voice-controlled cell phones. Ideally the interface would expand to handle more types of plans besides linear. There are several primary FRD-CSA systems involved architecturally, including the Universal Language (UniLang), Free Knowledge-Based System v2 (FreeKBS2), Verber, and Planning, Scheduling and Execution (PSE) systems.

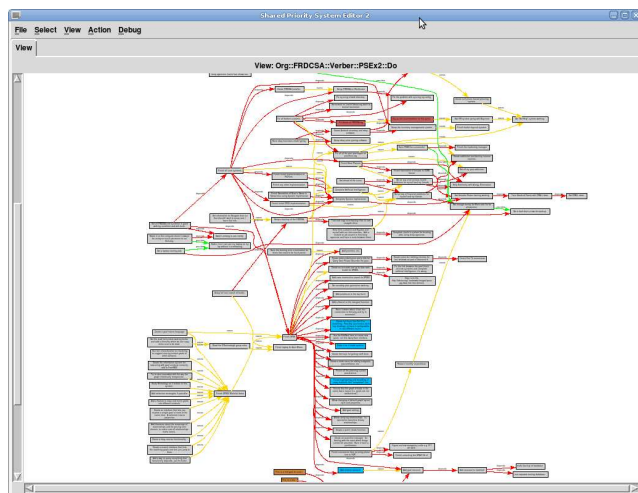


Figure 1: Shared Priority System Editor v2

System Architecture

Shared Priority System Editor

The name SPSE2 derives from Justin Coslor's concept of Priority Systems (Coslor 2008). Our goal system is straight-

forward, it has goals as nodes, unary predicates as conditions, and binary predicates as edges. Here is related meta-data from a sample goal in a planning domain. Suppose

```
<REL> is ( "entry-fn" "pse" "38" ):
( "asserter" <REL> "unknown" )
( "goal" <REL> )
( "has-NL" <REL> "ICAPS 2011 Paper" )
( "has-source" <REL>
  ( "entry-fn" "sayer-index" "806" ) )
( "depends" <REL> ( "entry-fn" "pse" "17" ) )
...
```

Goals are straightforwardly expressed in natural language, and may be marked with a unary predicate (such as **complete**, **showstopper**, **deleted**, **cancelled**, **ridiculous**, **obsoleted**, **rejected**, and **skipped**), or by a binary predicate (such as **depends**, **provides**, **eases** and **prefer**). SPSE2 is used to develop the planning domain, stored in a FreeKBS2 context (equivalent more or less to CYC's microtheories). The SPSE2 planning domain is converted to a domain usable by Verber, which then generates a PDDL domain. Implementation of more complex models of the semantics of processes is postponed until the basic system is completed. Eventually, goals that are ongoing or recurrent (i.e. do the laundry) should be modeled. For now, it only concerns whether a given goal has been completed. The SPSE2 GUI is being extended to become a knowledge editor, by adding several additional often related node types and predicate sets and background knowledge. There is also a domain for general knowledge modeling. The GUI has special functions for each domain. Domains under development include:

Alethic, Argumentation, CLEAR, Contexts, Critic, Deontic, DiscourseRepresentation, Doxastic, Genealogy, Intelligent-Agent, IntelligentTutoring, Inventory, Metamathematics, NetworkMapper, PICForm, Planning, POSI, Social-Networking, SuppositionalReasoner, Tactics, Temporal, Workflow

FreeKBS2

Knowledge is stored in FreeKBS2. FreeKBS2 can convert between several notations (KIF, Emacs and Perl Interlingua, CycL, etc) and will eventually have more back-ends besides the current Vampire-KIF back-end, enabling reasoning over higher-order and modal logics. Vampire provides first-order theorem proving with equality.

Verber

In order to generate a plan, SPSE2 translates its domain into one usable by Verber. Verber, named after the late Senior Chess Master Richard Verber, is basically a wrapper around various PDDL planners. Verber provides a set of Perl modules for building and interacting with PDDL domains (and hopefully other planning formalisms eventually) for calling various planners and parsing the results. It also houses a primitive knowledge engineering aspect for constructing Verb format planning domain libraries tailored to the specific domains required in the project, such as movement discipline, goal tracking and meal planning. The Verb format is

a very lightly extended PDDL, which allows importing other domains and problems and also a way to convert between PDDL's scalar time values and actual dates and times. Ideally Verber is capable of observation, learning the average and worst case durations of certain types of actions or events, and incorporating this into the plan development. Currently, there are several planners that are partially integrated, but so far LPG-TD is used primarily (Gerevini, Saetti, and Serina 2004).

UniLang

The UniLang system is an interprocess communication system for Perl "agents". UniLang is loosely termed a multiagent system, patterned off of the Open Agent Architecture, but without most of the Prolog-based communication capabilities (although it does have a trivial FreeKBS2-based knowledge interchange capability). Most agents, such as FreeKBS2, Verber and SPSE2 main and temporary agents, communicate with each other through the **Send** or **Query-Agent** functions.

PSE

PSE stands for Planning, Scheduling and Execution. It is one of the older FRDCSA planning systems. While the original design using object-oriented Perl code has been replaced entirely with the PDDL-based approach through Verber, there remains a significant collection of Emacs Lisp code within the namespace that works with the new Verber model. Before SPSE2, Emacs was the primary way of interacting with the planning system. Several interfaces were developed for manipulation of goals. Significant to note are the functions for rapidly asserting relations between goals. A stack-based interface, similar to an RPN scientific calculator, exists for pushing textual entities under the Emacs point onto the current stack, operating on the contents of the current stack and ring of stacks, and asserting into FreeKBS2. One such function will return the entry ID of the goal corresponding to the text under the Emacs point. This interface is still very much under development for the Natural Language Understanding (NLU) system and the Knowledge-Max editor (KMax) systems, allowing semantic markup of text and consequent rapid manipulation of symbolic, textual knowledge.

System Interfaces

Calendar Synchronization

In order to improve the utility of Verber, a calendar synchronization feature has been developed. This allows synchronization with ICS Calendars or Google Calendars, the ICS files for which are obtained and then translated into the SPSE2 domain representation. In order to schedule an event occurring at a certain time, it is translated from the date and time information into the offset and scale of the scalar planning time value. The event start date is marked using PDDL timed initial literals.

```
(at 1.00 (possible <EVENT>))
(at 2.00 (not (possible <EVENT>)))
```

Additionally the planning domain includes the following precondition for the durative action **Complete**.

```
(over all
  (or
    (not (has-time-constraints ?e1))
    (possible ?e1)))
```

Here is an example of the time constraints that are asserted by the DateTime interface.

```
("end-date"
  ("entry-fn" "pse" "38")
  "TZID=America/Chicago:20101129T120000")
```

Notification Manager

A Perl/Tk-based notification manager patterned off of the Android notification manager has been completed partially. Ideally, it would be able to synchronize with the Android notification manager. Certain tasks that regularly occur are regularly added to the SPSE2 planning context by a cron job. The custom cron-like format includes information on how much warning time should predate the beginning of the possibility of completing a task.

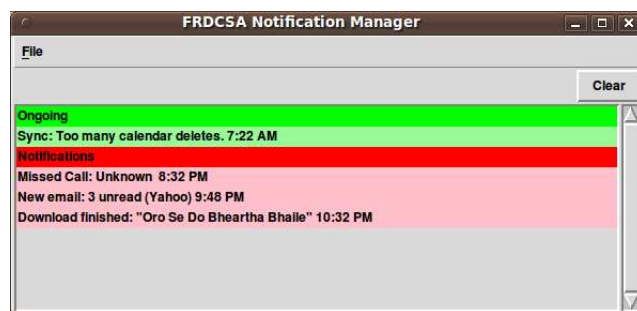


Figure 2: Notification Manager

Deontic and Teleologic Logics

It should be possible for users to place certain ethical constraints on possible actions, in order to provide their agents with a moral decision support system. In support of this, work is ongoing to formalize various systems of morality in terms of deontic logics, and to provide an evaluation function which evaluates individual actions and plans against various selected moral systems. Therefore, the user simply declares the moral systems to which they agree, and the system then evaluates the actions. To illustrate a glaring but simple example, let us consider the Ten Commandments, specifically the obligation that "Thou shalt not kill". If our computational semantics lacked understanding of general prohibitions, perhaps this could be manually translated to "Someone or something murdered". This is then converted to a logic form <LF> and it asserts:

```
(implies <LF> (rule-activated <RULE-NUMBER>))
```

It then determines whether this rule is activated using semantic textual entailment recognition. The recognizer would load the logic of the plan and all enumerable consequences,

convert to logic form, instantiate the variables with their valuations, add contingent background knowledge, and query: (rule-activated ?X), relaxing the lexical constraints until a rule is activated. Of course, in practice this will be more complicated (Balduccini, Baral, and Lierler 2007). Besides considering general obligations and prohibitions (deontics), it should also be useful to consider a means-ends analysis (teleologics).

World State Comparison and Value Systems

An intended capability is to enable reasoning intelligently about the value of a given plan or outcome. In trying to evaluate various actions to determine which are morally superior according to various moral theories, it begs the question of which world-states and histories are preferable. To evaluate this, a general comparison function needs to be developed. To reduce the evaluation to a total order would be somewhat dualistic - but ultimately the user should be able to choose preferable from among possible worlds - or rather provide rules which in some sense order them.

Android Bluetooth Headset-Based Interactive Execution Monitor

In order to provide a usable interface to the plans, an application for the Google Android mobile operating system is being developed for recording new goals, walking the user through generated plans, and initiating replanning. Currently, a client-server system has been implemented that communicates between the phone and server using XMLRPC. This functionality actually resides within UniLang. A limited set of voice commands has been developed, enabling the phone to for instance answer factoid-based questions using our agent wrapping the open source OpenEphyra question answering system. Here is an illustrated use case for this kind of question answering. The user initiates the speech interface by pressing a certain combination of buttons on their Bluetooth headset, currently only the media button. Recognition occurs using Google's voice recognition API, and results are then sent via XMLRPC to the Android-FRDCSA-Server UniLang agent. The command is processed by a VoiceCommand module, which in this use case sends a query to the QUAC Question Answering system, before ultimately returning the message.

Previously, a text-based interactive execution monitor was deployed which used the Sphinx2 numeric domain and TTS to query the user as to the status of the completed plan. Thus, it is not hard to imagine adapting the current text-based interactive execution manager to work on the Android phone. Progress was delayed for months while the development phone was broken, but thankfully another unit has been procured recently. Initiating communication to the phone, on the other hand, while more difficult, may be accomplished either through registering the phone with a dynamic-DNS provider, or by polling from the phone to the server. Voice commands for controlling the process of planning, replanning and the interactive execution monitor will be developed. Development of the interactive execution monitor, especially one that incorporates conditional planning for various possible interaction situations (such as whether the user

is in the presence of the phone or the computer), will be more difficult.

Location Logic

Location Logic is a system for inferencing with the semantics of logic related to the users' physical locations as reported by the GPS (or in the multiagent case using location services like Google Latitude), as contrasted to their way-point data. It works by asserting theorems observed regarding the GPS tracks, such as whether certain waypoints are being approached, visited, or departed. Location Logic interfaces naturally with the SPSE2 system. Locational restrictions will be extracted from goals using named entity recognition and computational semantics. Eventually, multiagent, distributed planning solvers should be integrated which are able to resolve distributed resource problems.

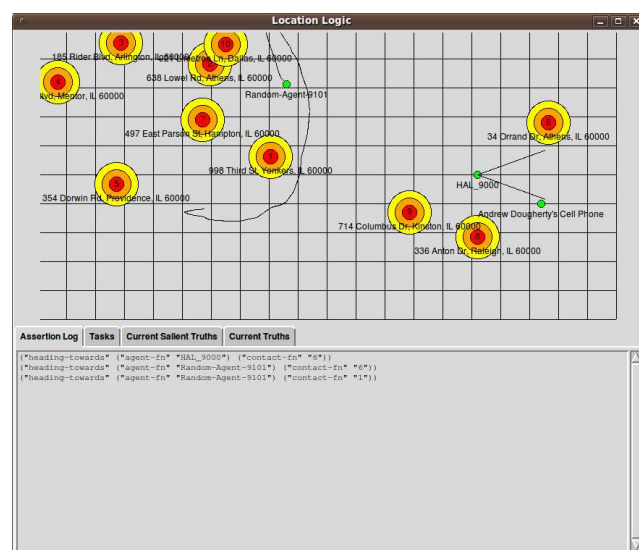


Figure 3: Location Logic Prototype

Here is an example use case. Suppose both Alice and Bob are using mobile phones having the hands-free voice-controlled task list and Location Logic systems installed. Both Alice and Bob are wearing Bluetooth headsets.

So suppose Alice is out of milk, but she doesn't know this right away. (Perhaps her roommate drank the last of it, and did not or could not tell her yet). Meanwhile, Bob is out running errands of his own. Alice wakes up and comes down the stairs and goes to get some cereal. She pours the bowl of cereal and then looks in her refrigerator. She realizes that the milk is empty and has not been thrown out. Disappointed, she taps a button on her Bluetooth headset. It responds "Yes?". Alice says: "pantry remove item milk", to which her headset responds, "Confirm remove milk from pantry inventory?". Alice says "yes" and the headset responds "Milk removed".

When Alice tells the system that she was out of milk - the system realizes according to various considerations that she should buy some more milk. It therefore automatically added the goal "get milk". Because Alice and Bob are

friends, they have already told their Shared Priority Systems that they can collaborate on several matters, one of which is food inventory. Perhaps Alice's planning agent sent a broadcast to all of her friends with which she shares this particular kind of information. It was stated simply as the goal that there should be a fresh jug of milk in Alice's refrigerator. The various agents therefore add this goal to their own planning systems. Collectively, at the next planning or replanning cycle, using a distributed planning formalism, they agree on the details of the least costly and most preferable plan that incorporates Alice's request. Bob's planning agent and ethical analyzer agree to propose to Bob the task of picking up and dropping off milk to Alice's house, because he was approaching a grocery store, didn't have other things that couldn't be slightly postponed, and his adjustable autonomy settings allowed for the proposal. Bob and perhaps Alice get a message suggesting that this action be taken. Once both have agreed, it is added to the planning system, even including so much as to tell Bob's cell phone agent which type of milk Alice would like. Bob then purchases the milk and then swings by Alice's place on his way to work. For now, assume Alice pays him cash - but in the future an automated distributed loan/payment system will be used.

The usefulness of the milk provision scenario is debatable, but certainly in the general case such efficient team-based collaboration would be very desirable, as for instance in the case of ride-sharing, shared task management, and so on.

Here is an example Location Logic rule.

```
(implies
  (and
    (leaving ?AGENT ?LOCATION)
    (isa ?LOCATION movie-theatre)
    (has-performed-action ?AGENT
      "silence cell phone at movie theatres")
    ;: (> (sitting-still) (minute 1))
  )
  (perform-action "add-to-pending-tasks"
    "unsilence cell phone \
      when leaving movie theatres"))
```

Federated Transportation Planning

A federated transportation planning option, based on the FRDSCA BusRoute planner¹ has been integrated with Verber. However, with the advent of Google's public transportation planner it will be replaced. This interface should naturally understand waypoints, and be able to generate queries to a public or private transportation routing system in order to populate the planning domain with the timing constraints. Ideally services transmitting actual bus positions and timing could be integrated - and replanning initiated as needed.

Systems Using SPSE2/Verber/PSE

POSI Collaboration

The planning system and interactive execution monitor have several applications. One related project using the tools of

the FRDSCA is the POSI Open Source Initiative (POSI)². It is a project for representing the goals, interests and abilities of its users, as distilled from their writings and volunteered information. The idea is to establish necessary and sufficient information to form dynamic multiagent teams to solve problems that are shared between multiple persons. Priority Systems are essentially networks of goals and constraints on these goals. Future versions of SPSE will be able to simultaneously edit the same Shared Priority Systems. Various models of shared editing are under consideration. Ideally, when someone marks a goal as completed, the other SPSE agents should display such changes. An analogy can be made to a real-time strategy game. An essential feature of POSI is identifying when different users have specified the same or related goals. This is accomplished chiefly using the nascent technology of Textual Entailment Recognition, as well as through breaking down individual goals into more clearly defined or achievable steps. As such, the Goals, Interests and Abilities (GIAs) of the users are modeled using ontological tools. SPSE2 itself soon will have the ability to edit domains besides temporal planning, including different ontologies such as the POSI ontologies.

Akahige Medical System

Another intended application is the Akahige Medical System³. One capability of the Android phone-based general purpose help system under development is to launch a medical help system in the event that any conspicuous symptoms occur or at regularly scheduled times. Symptoms may be run through a standard or custom medical diagnostic program (such as Diagnosaurus or the planned Akahige Model-Based Diagnostic and Fault Localization software). The result of the diagnostic procedure may require an emergency response. In such a case, instructions for the given emergency or situation will exist within the system and the user will be guided by the interactive execution monitor to complete these tasks, and to refer to documentation as needed.

Gourmet Meal Planner

Another area of significant interest is meal planning. The FRDSCA project includes a meal planner, called Gourmet⁴. It will interface with our inventory management component for pantry management. The SOAR archive of around 150,000 meal-master recipes is being integrated. Work is progressing on developing a foodstuffs ontology, upon which to map the ingredients and intermediate foodstuffs of recipes. Mapping to nutrient databases such as the SR23 and also to product databases such as UPCdatabase.com, in addition to formalizing the recipe steps into discrete planning operations, will provide all the resources required to achieve interactive recipe execution through the Android interface. The formalization of recipe steps will hopefully be attempted by training CMU's StackedFrameParser on the CURD/MILK dataset of annotated recipes. They use a medium grained language called CURD which expresses

²<http://intranet.posi.frdcsa.org>

³<http://frdcsa.org/frdcsa/internal/akahige>

⁴<http://frdcsa.org/frdcsa/internal/gourmet>

¹<http://frdcsa.org/frdcsa/internal/busroute>

a set of abstracted operations on food and intermediate foodstuffs (Tasse and Smith 2008). Then arbitrary English recipes may be attempted. As some other open source meal planners have well established user bases, this automatic annotation could be performed and cached by a web server (as the StackedFrameParser requires 8GB of RAM), and could be augmented and trained by user corrections given a sufficient GUI and upload mechanism.

Paperless-Office System

Another FRDCSA project that will make use of the planning technology is the Paperless-Office⁵. With this system, which already functions to scan, OCR, search and edit documents, it will be possible to specify workflows - such as certain documents requiring to be filled out and sent by certain times.

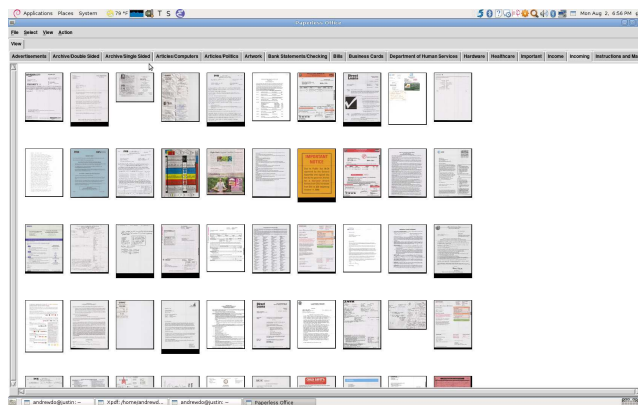


Figure 4: Paperless-Office

SystemX Intelligent Tutoring System

Work has been ongoing on Arbitrary Document Understanding. Being able to represent the argument structure of a text, and also the facts and relations of the text, enables the generation of temporal plans for teaching subjects at various granularities based on various learning objectives and timing constraints. A prototype system called Study has been developed⁶. SystemX will use the SPSE2 or its derivatives.

Conclusions and Future Work

Usable Perl/Tk, Emacs and Android systems for editing goals, adding temporal constraints, generating plans, and walking the user through them are close to completion. The chief obstacle to the release of this work has been its thorough dependence on the FRDCSA, a large set of heavily interconnected yet unreleased software. A proper requirements engineering process should be initiated (as the project management for SPSE2 is currently performed within SPSE2 itself - a chicken and egg problem). The completion of the Android interactive execution monitor is now

possible following the recent acquisition of a working Android phone. As well, some pernicious bugs affect the FRDCSA notification manager.

There are several other related topics that bear mentioning but that unfortunately time constraints have precluded. This includes plan libraries (including some other previously developed PDDL domains for personal tasks involving several other predicates), ontological modeling, scripts and automatic execution of tasks, house rules, textual entailment, computational semantics via logic forms, static domain analysis, reasoning with the consequences of failing to complete certain tasks, automated goal analysis, an ontology of planning systems and their capabilities, plan cycles, and incorporating conformant and other plan types into the interactive execution monitor. Here are some important areas for future work.

Automatic PDDL Domain Construction Via Computational Semantics

Although currently goals are simply evaluated in a boolean context, a more detailed interpretation of the semantics of the goals is planned. Ideally, it would convert the natural language contents of goals and other node types into a logical semantic representation, such as logic forms (LFs), and from the LFs construct the PDDL domains and problems.

Suppositional Reasoner

The suppositional reasoner seeks to incorporate more positional evaluation and analysis of domain invariants, domain specific knowledge and so on, into the planning process. It ideally would function as a plan development and critiquing interface for real-life problems, hopefully allowing the evaluation of any decision making process. It is being calibrated on domains like Chess and Go that have a substantial literature that may be formalized and in which feedback on efficacy is relatively short. From a search point of view, it is not very interesting, for starters some standard search algorithms should be implemented. But attempting to hybridize the search with positional information is interesting. Annotated chess games and their logic-based formalizations are being used to develop knowledge that may be tested and applied to the situation. The Sayer system provides a method for a kind of theorem proving over these domains and to perform Natural Language Understanding (NLU) within the project.

Temporal Conformant Planning

One desired capability is temporal conformant planning, and tools for exploring the set of actions possible to the agent at each state, and reasoning with the consequences of various choices at that point. This naturally reflects back to the suppositional reasoner and the valuation system. This is perhaps related to Conditional Temporal Planning (Tsamardinos, Vidal, and Pollack 2003). Tools that implement these techniques should be located and integrated.

⁵<http://freshmeat.net/paperless-office>

⁶<http://frdcsa.org/frdcsa/internal/study>

Acknowledgments

I am grateful especially to Aloysius Flori and my immediate family who have been supportive of my work; to Justin Coslor who has provided the inspiration for countless systems and programs and has developed a formal theory of context; to the CMU community for its helpfulness; to Jim Oberweis and the late Richard Verber for their support and mentorship with chess; and to the many persons who showed me kindnesses when they were under no obligation to do so.

I wish to thank all of the researchers who have or will have provided references and pointers, those who have made progress in this field, and in particular those who have released software, documentation and papers under freely redistributable or preferably FOSS licenses, such as the authors of LPG-TD. Such licenses make the work of an open source systems integrator much easier.

References

- Balduccini, M.; Baral, C.; and Lierler, Y. 2007. Knowledge representation and question answering. In Harmelen, V.; Lifschitz; and Porter., eds., *Handbook of Knowledge Representation*. chapter 21.
- Chaitin, G. J. 1974. Information-theoretic limitations of formal systems. *Journal of the ACM* 21:403–424.
- Coslor, J. M. 2008. *Possibility Thinking*. Lulu.com.
- Gerevini, A.; Saetti, A.; and Serina, I. 2004. LPG-TD: a fully automated planner for PDDL2.2 domains. *14th Int. Conference on Automated Planning and Scheduling (ICAPS-04)*.
- Martin, C. 2006. Executive skills: How to improve your ability to focus.
- Tasse, D., and Smith, N. A. 2008. SOUR CREAM: Toward semantic processing of recipes. Technical Report 5, Carnegie Mellon University, Pittsburgh, PA.
- Tsamardinos, I.; Vidal, T.; and Pollack, M. 2003. CTP: A new constraint-based formalism for conditional, temporal planning.
- Turing, A. 1939. Systems of logic based on ordinals. In *Proc. London Math. Soc.*, 161–228.